

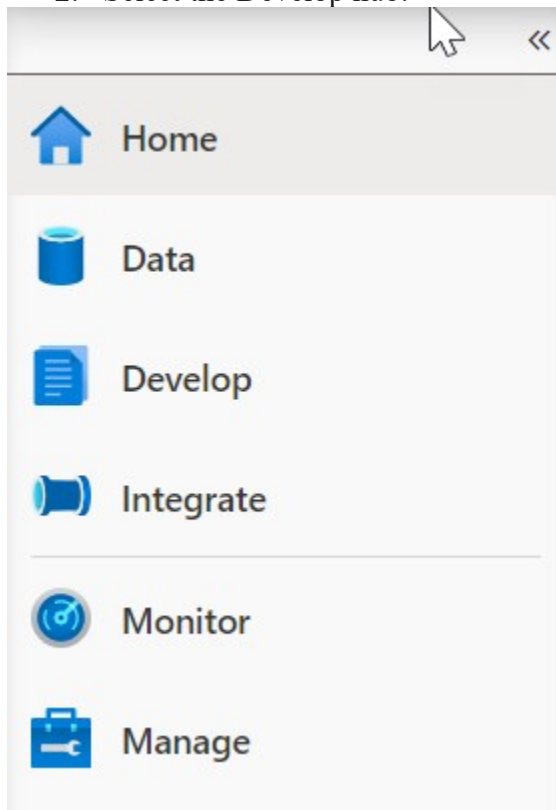
Work with windowing functions

Note:

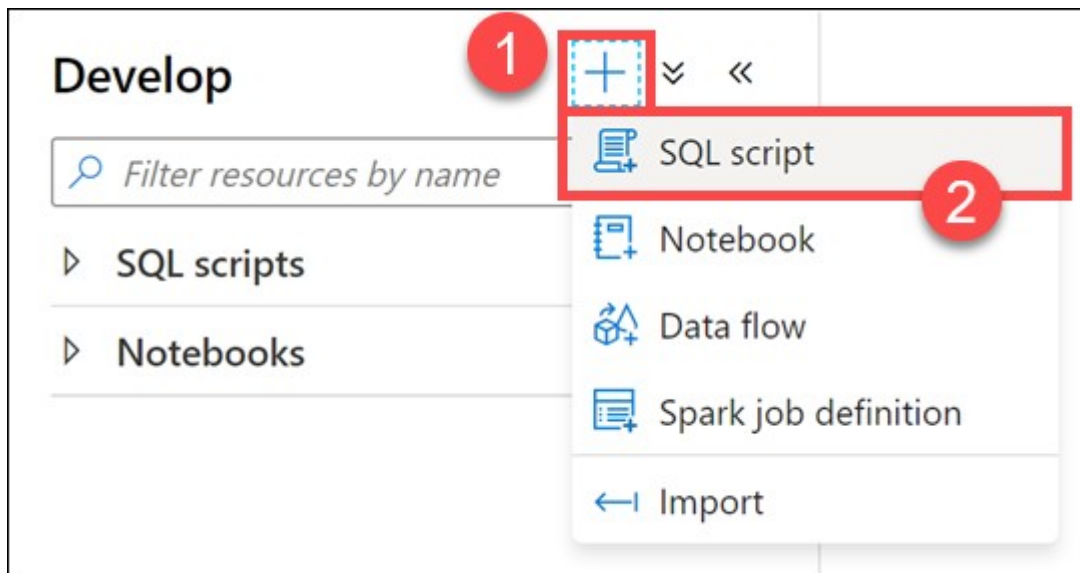
You are not required to complete the processes, tasks, activities, or steps presented in this example. The various samples provided are for illustrative purposes only and it's likely that if you try this out you will encounter issues in your system.

One of the key components of window functions is the OVER clause. This clause determines the partitioning and ordering of a rowset before the associated window function is applied. That is, the OVER clause defines a window or user-specified set of rows within a query result set. A window function then computes a value for each row in the window. You can use the OVER clause with functions to compute aggregated values such as moving averages, cumulative aggregates, running totals, or a top N per group results.

1. Open Synapse Studio (<https://web.azuresynapse.net/>).
2. Select the Develop hub.



3. From the Develop menu, select the + button (1) and choose SQL Script (2) from the context menu.



4. In the toolbar menu, connect to the SQL Pool database to execute the query.



5. In the query window, replace the script with the following to use the OVER clause with data from the wwi_security.Sale table:

```
SELECT
ROW_NUMBER() OVER(PARTITION BY Region ORDER BY Quantity DESC) AS "Row Number",
Product,
Quantity,
Region
FROM wwi_security.Sale
WHERE Quantity <> 0
ORDER BY Region;
```

6. Select Run from the toolbar menu to execute the SQL command.



When we use PARTITION BY with the OVER clause (1), we divide the query result set into partitions. The window function is applied to each partition separately and computation restarts for each partition.

1
2
3
4
5
6
7
8

1	SELECT			
2	ROW_NUMBER()	OVER(PARTITION BY Region	ORDER BY Quantity DESC)	AS "Row Number",
3	Product,			
4	Quantity,			
5	Region			
6	FROM wwi_security.Sale			
7	WHERE Quantity <> 0			
8	ORDER BY Region;			

Results
Messages

View
Table
Chart
Export results

Search

Row Number	Product	Quantity	Region
16	Wheel	5	Far West
17	Video Games	5	Far West
18	Video Games	5	Far West
19	Wireless headphones	3	Far West
20	Wireless headphones	3	Far West
21	Key Chain	2	Far West
22	Key Chain	2	Far West
1	Sunglasses	52	South East
2	Beach hat	52	South East

The script we executed uses the OVER clause with ROW_NUMBER function (1) to display a row number for each row within a partition. The partition in our case is the Region column. The ORDER BY clause (2) specified in the OVER clause orders the rows in each partition by the column Quantity. The ORDER BY clause in the SELECT statement determines the order in which the entire query result set is returned.

Scroll down in the results view until the Row Number count (3) starts over with a different region (4). Since the partition is set to Region, the ROW_NUMBER resets when the region changes. Essentially, we have partitioned by region and have a result set identified by the number of rows in that region.

Aggregate functions

Now use aggregate functions with our window by expanding on our query that uses the OVER clause.

1. In the query window, replace the script with the following to add aggregate functions:

3
4
5
6
7
8
9
10
11
12
13

```
SELECT
ROW_NUMBER() OVER(PARTITION BY Region ORDER BY Quantity DESC) AS "Row Number",
Product,
Quantity,
SUM(Quantity) OVER(PARTITION BY Region) AS Total,
AVG(Quantity) OVER(PARTITION BY Region) AS Avg,
COUNT(Quantity) OVER(PARTITION BY Region) AS Count,
MIN(Quantity) OVER(PARTITION BY Region) AS Min,
MAX(Quantity) OVER(PARTITION BY Region) AS Max,
Region
FROM wwi_security.Sale
WHERE Quantity <> 0
ORDER BY Region;
```

2. Select Run from the toolbar menu to execute the SQL command.



In our query, we added the SUM, AVG, COUNT, MIN, and MAX aggregate functions. Using the OVER clause is more efficient than using subqueries.

5	SUM(Quantity) OVER(PARTITION BY Region) AS Total,
6	AVG(Quantity) OVER(PARTITION BY Region) AS Avg,
7	COUNT(Quantity) OVER(PARTITION BY Region) AS Count,
8	MIN(Quantity) OVER(PARTITION BY Region) AS Min,
9	MAX(Quantity) OVER(PARTITION BY Region) AS Max,
10	Region
11	FROM wwi_security.Sale
12	WHERE Quantity <> 0
13	ORDER BY Region;

Results Messages

View Table Chart Export results

Search

Row Number	Product	Quantity	Total	Avg	Count	Min	Max	Region
17	Video Games	5	328	14	22	2	32	Far West
18	Video Games	5	328	14	22	2	32	Far West
19	Wireless headphones	3	328	14	22	2	32	Far West
20	Wireless headphones	3	328	14	22	2	32	Far West
21	Key Chain	2	328	14	22	2	32	Far West
22	Key Chain	2	328	14	22	2	32	Far West
1	Sunglasses	52	962	32	30	5	52	South East
2	Beach hat	52	962	32	30	5	52	South East
3	Beach hat	52	962	32	30	5	52	South East
4	Sunglasses	52	962	32	30	5	52	South East

Analytic functions

Analytic functions calculate an aggregate value based on a group of rows. Unlike aggregate functions, however, analytic functions can return multiple rows for each group. Use analytic functions to compute moving averages, running totals, percentages, or top-N results within a group.

Tailwind Traders has book sales data they import from their online store and wish to compute percentages of book downloads by category.

To do this, you decide to build window functions that use the PERCENTILE_CONT and PERCENTILE_DISC functions.

1. In the query window, replace the script with the following to add aggregate functions:

```
-- PERCENTILE_CONT, PERCENTILE_DISC
SELECT DISTINCT c.Category
```

1
2
3
4
5
6
7
8
9
10
11
12

```

,PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY bc.Downloads)
    OVER (PARTITION BY Category) AS MedianCont
,PERCENTILE_DISC(0.5) WITHIN GROUP (ORDER BY bc.Downloads)
    OVER (PARTITION BY Category) AS MedianDisc
FROM dbo.Category AS c
INNER JOIN dbo.BookList AS bl
    ON bl.CategoryID = c.ID
INNER JOIN dbo.BookConsumption AS bc
    ON bc.BookID = bl.ID
ORDER BY Category

```

2. Select Run from the toolbar menu to execute the SQL command.

The screenshot shows the SQL Server Enterprise Manager interface. The top toolbar has a 'Run' button highlighted with a red box and a red circle labeled '1'. The query editor contains the following SQL code:

```

1  -- PERCENTILE_CONT, PERCENTILE_DISC
2  SELECT DISTINCT c.Category
3  ,PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY bc.Downloads)
4  ,PERCENTILE_DISC(0.5) WITHIN GROUP (ORDER BY bc.Downloads)
5  FROM dbo.Category AS c
6  INNER JOIN dbo.BookList AS bl
7  ON bl.CategoryID = c.ID
8  INNER JOIN dbo.BookConsumption AS bc
9  ON bc.BookID = bl.ID
10 ORDER BY Category

```

Red boxes and arrows highlight specific parts of the query: a red box labeled '2' around the window functions, a red box labeled '3' around the 'WITHIN GROUP' clause, and a red arrow labeled '4' pointing to the 'ORDER BY' clause.

Below the query editor, the 'Results' tab is selected. The 'View' dropdown is set to 'Table'. The results are displayed in a table with the following data:

Category	MedianCont	MedianDisc
Favorites	304.75	231.833333333333
Fiction	1091.58333333333	935
Romance	78.5833333333333	59.1666666666667
Sci-fi	54.75	44.3333333333333
Trending Now	187	153.333333333333

In this query, we use PERCENTILE_CONT (1) and PERCENTILE_DISC (2) to find the median number of downloads in each book category. These functions may not return the same value. PERCENTILE_CONT interpolates the appropriate value, which may or may not exist in the data set, while PERCENTILE_DISC always returns an actual value from the set. To explain further, PERCENTILE_DISC computes a specific percentile for sorted values in an entire rowset or within a rowset's distinct partitions.

The 0.5 value passed to the percentile functions (1 & 2) computes the 50th percentile, or the median, of the downloads.

The WITHIN GROUP expression (3) specifies a list of values to sort and compute the percentile over. Only one ORDER BY expression is allowed, and the default sort order is ascending.

The OVER clause (4) divides the FROM clause's result set into partitions, in this case, Category. The percentile function is applied to these partitions.

3. In the query window, replace the script with the following to use the LAG analytic function:

```
--LAG Function
SELECT ProductId,
       [Hour] AS SalesHour,
       TotalAmount AS CurrentSalesTotal,
       LAG(TotalAmount, 1,0) OVER (ORDER BY [Hour]) AS PreviousSalesTotal,
       TotalAmount - LAG(TotalAmount,1,0) OVER (ORDER BY [Hour]) AS Diff
FROM [wwi_perf].[Sale_Index]
WHERE ProductId = 3848 AND [Hour] BETWEEN 8 AND 20;
```

Tailwind Traders wants to compare the sales totals for a product over an hourly basis over time, showing the difference in value.

To accomplish this, you use the LAG analytic function. This function accesses data from a previous row in the same result set without the use of a self-join. LAG provides access to a row at a given physical offset that comes before the current row. We use this analytic function to compare values in the current row with values in a previous row.

<pre> 1 --LAG Function 2 SELECT ProductId, 3 [Hour] AS SalesHour, 4 TotalAmount AS CurrentSalesTotal, 5 LAG(TotalAmount, 1,0) OVER (ORDER BY [Hour]) AS PreviousSalesTotal, 6 TotalAmount - LAG(TotalAmount,1,0) OVER (ORDER BY [Hour]) AS Diff 7 FROM [wwi_perf].[Sale_Index] 8 WHERE ProductId = 3848 AND [Hour] BETWEEN 8 AND 20; </pre>																																																											
<div>Results</div> <div>Messages</div> <div>View</div> <div>Table</div> <div>Chart</div> <div>Export results</div>																																																											
<div>Search</div> <table> <tr> <th>ProductId</th><th>SalesHour</th><th>CurrentSalesTotal</th><th>PreviousSalesTotal</th><th>Diff</th></tr> <tr> <td>3848</td><td>8</td><td>27.71</td><td>0.00</td><td>27.71</td></tr> <tr> <td>3848</td><td>8</td><td>55.42</td><td>27.71</td><td>27.71</td></tr> <tr> <td>3848</td><td>8</td><td>27.71</td><td>55.42</td><td>-27.71</td></tr> <tr> <td>3848</td><td>8</td><td>110.84</td><td>27.71</td><td>83.13</td></tr> <tr> <td>3848</td><td>8</td><td>55.42</td><td>110.84</td><td>-55.42</td></tr> <tr> <td>3848</td><td>8</td><td>110.84</td><td>55.42</td><td>55.42</td></tr> <tr> <td>3848</td><td>8</td><td>110.84</td><td>110.84</td><td>0.00</td></tr> <tr> <td>3848</td><td>8</td><td>55.42</td><td>110.84</td><td>-55.42</td></tr> <tr> <td>3848</td><td>8</td><td>83.13</td><td>55.42</td><td>27.71</td></tr> <tr> <td>3848</td><td>8</td><td>27.71</td><td>83.13</td><td>-55.42</td></tr> </table>					ProductId	SalesHour	CurrentSalesTotal	PreviousSalesTotal	Diff	3848	8	27.71	0.00	27.71	3848	8	55.42	27.71	27.71	3848	8	27.71	55.42	-27.71	3848	8	110.84	27.71	83.13	3848	8	55.42	110.84	-55.42	3848	8	110.84	55.42	55.42	3848	8	110.84	110.84	0.00	3848	8	55.42	110.84	-55.42	3848	8	83.13	55.42	27.71	3848	8	27.71	83.13	-55.42
ProductId	SalesHour	CurrentSalesTotal	PreviousSalesTotal	Diff																																																							
3848	8	27.71	0.00	27.71																																																							
3848	8	55.42	27.71	27.71																																																							
3848	8	27.71	55.42	-27.71																																																							
3848	8	110.84	27.71	83.13																																																							
3848	8	55.42	110.84	-55.42																																																							
3848	8	110.84	55.42	55.42																																																							
3848	8	110.84	110.84	0.00																																																							
3848	8	55.42	110.84	-55.42																																																							
3848	8	83.13	55.42	27.71																																																							
3848	8	27.71	83.13	-55.42																																																							

In this query, we use the LAG function (1) to return the difference in sales (2) for a specific product over peak sales hours (8-20). We also calculate the difference in sales from one row to the next (3). Notice that because there is no lag value available for the first row, the default of zero (0) is returned.

ROWS clause

The ROWS and RANGE clauses further limit the rows within the partition by specifying start and end points within the partition. This is done by specifying a range of rows with respect to the current row either by logical association or physical association. Physical association is achieved by using the ROWS clause.

Tailwind Traders wants to find the books that have the fewest downloads by country, while displaying the total number of downloads for each book within each country in ascending order.

To achieve this, you use ROWS in combination with UNBOUNDED PRECEDING to limit the rows within the Country partition, specifying that the window start with the first row of the partition.

1. In the query window, replace the script with the following to add aggregate functions:

```
-- ROWS UNBOUNDED PRECEDING
SELECT DISTINCT bc.Country, b.Title AS Book, bc.Downloads
    ,FIRST_VALUE(b.Title) OVER (PARTITION BY Country
        ORDER BY Downloads ASC ROWS UNBOUNDED PRECEDING) AS FewestDownloads
FROM dbo.BookConsumption AS bc
INNER JOIN dbo.Books AS b
    ON b.ID = bc.BookID
ORDER BY Country, Downloads
```

2. Select Run from the toolbar menu to execute the SQL command.



1

-- ROWS UNBOUNDED PRECEDING

2

SELECT DISTINCT bc.Country, b.Title AS Book, bc.Downloads

3

,FIRST_VALUE(b.Title) OVER (PARTITION BY Country

4

ORDER BY Downloads ASC ROWS UNBOUNDED PRECEDING) AS FewestDownloads

5

FROM dbo.BookConsumption AS bc

6

INNER JOIN dbo.Books AS b

7

ON b.ID = bc.BookID

8

ORDER BY Country, Downloads

Results

Messages

View

Table

Chart

Export results

Search

Country	Book	Downloads	FewestDownloads
Germany	LOVE LIKE THAT (The Romance Ch...	1233.0000000000007	Harry Potter - The Ultimate Quiz Book
Germany	A Reaper at the Gates	1418	Harry Potter - The Ultimate Quiz Book
Germany	Harry Potter - The Ultimate Book...	1565.6666666666667	Harry Potter - The Ultimate Quiz Book
Germany	Harry Potter - The Amazing Quiz ...	1721.6666666666667	Harry Potter - The Ultimate Quiz Book
Germany	101 Amazing Harry Potter Facts	2410.3333333333333	Harry Potter - The Ultimate Quiz Book
Germany	Harry Potter - The Ultimate Quiz	2812.6666666666667	Harry Potter - The Ultimate Quiz Book
Sweden	Burn for Me	1.1999999999999999	Burn for Me
Sweden	Diamond Fire	6	Burn for Me
Sweden	Game Of Thrones The Quiz Book...	8.1999999999999999	Burn for Me
Sweden	Harry Potter - The Complete Qui...	11	Burn for Me
Sweden	The World of Ice and Fire: The U...	15.6	Burn for Me

In this query, we use the FIRST_VALUE analytic function to retrieve the book title with the fewest downloads, as indicated by the ROWS UNBOUNDED PRECEDING clause over the Country partition (1). The UNBOUNDED PRECEDING option set the window start to the first row of the partition, giving us the title of the book with the fewest downloads for the country within the partition.

In the result set, we can scroll through the list that of books by country, sorted by number of downloads in ascending order.